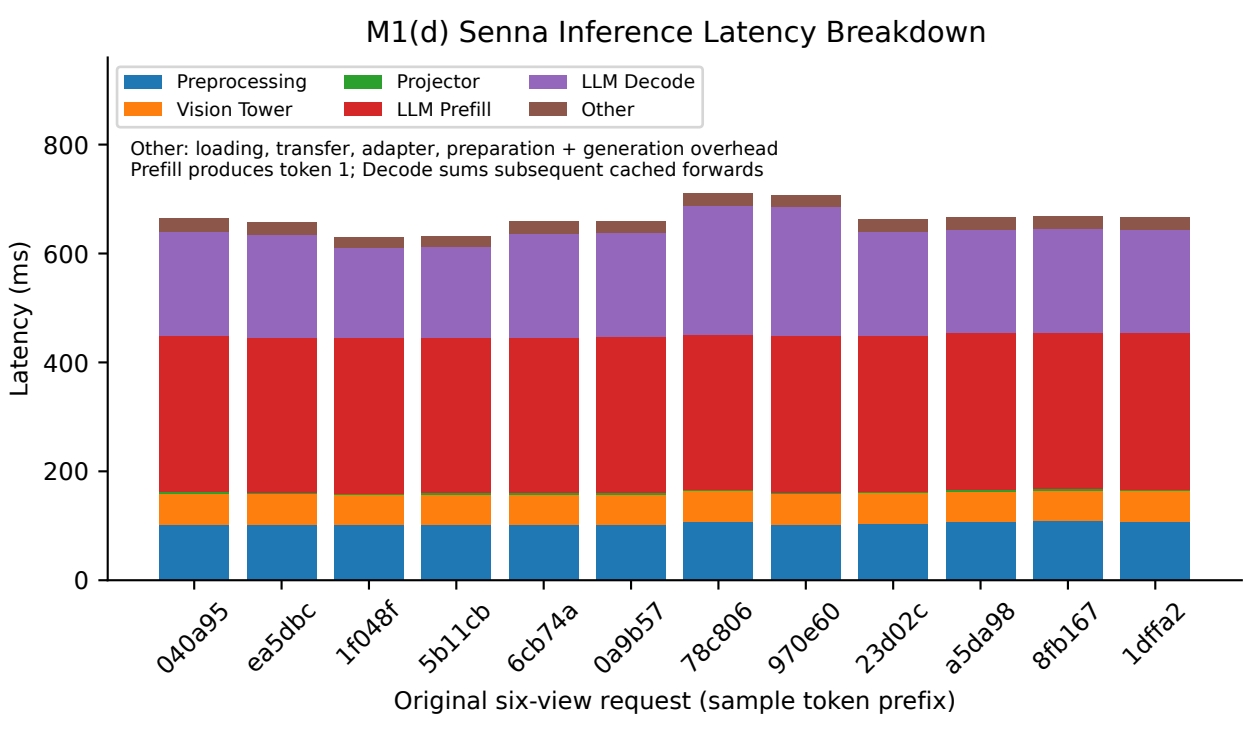
M1(d) Senna Inference Latency Breakdown
Preprocessing
Vision Tower
Projector
LLM Prefill
LLM Decode
Other
Other: loading, transfer, adapter, preparation + generation overhead
Prefill produces token 1; Decode sums subsequent cached forwards
Latency (ms)
0
200
400
600
800
040a95
ea5dbc
1f048f
5b11cb
6cb74a
0a9b57
78c806
970e60
23d02c
a5da98
8fb167
1dffa2
Original six-view request (sample token prefix)